

Shaurya Gulati

i.shauryagulati@gmail.com · [linkedin.com/in/shauryagulati](https://www.linkedin.com/in/shauryagulati) · github.com/Shaurayagulati · shauryagulati.me

EDUCATION

Carnegie Mellon University — Master's in Information Systems Management, Business Intelligence & Data Analytics (Aug 2024 – Dec 2025)

Coursework: Computational Data Science, Machine Learning, Applications of NLP and LLMs, Distributed Systems, Applied AI.

Chandigarh University — B.E. Computer Science, specialization in Artificial Intelligence & Machine Learning (2018 – 2022)

Undergraduate researcher in sentiment analysis and computer vision; published at international conferences.

EXPERIENCE

Lunon — *Founding AI Engineer*

Feb 2026 – May 2026

- Built the 5-phase core AI pipeline for automated commercial due diligence (private equity): a multi-agent system with quality gates and automated correction loops.
- Built the retrieval stack: document ingestion, hybrid search (pgvector + lexical with RRF fusion), query decomposition, and cross-encoder reranking.
- Built the analyst copilot: a ~20-tool agent with streaming chat, a propose-then-confirm flow for edits, and citation verification against the source corpus.
- Traced and fixed a quality-gate failure loop blocking a client demo; added warn-mode so runs completed instead of stalling.
- Set up Langfuse observability for tracing and cost/latency monitoring across the agent pipeline.

Carnegie Mellon University — *AI Research Assistant*

Jun 2025 – Dec 2025

- Engineered a hybrid RAG system across 1,000+ regulatory documents: dense retrieval (E5 embeddings, ChromaDB) fused with sparse BM25, re-ranked with a BGE cross-encoder.
- Built a PDF ingestion pipeline with page parsing, header detection, and overlapping chunking, with data-lineage tracking for accurate vector-store population.
- Evaluated across 50+ query sets: 94% accuracy with grounded citations.

Carnegie Mellon University, Heinz College — *Graduate Teaching Assistant, Fundamentals of Operationalizing AI*

Aug 2025 – Oct 2025

- Ran recitations and Q&A on deploying and scaling AI systems in practice; guided students through implementation challenges and graded coursework.

YMGrad — *Software Development Engineer*

Jun 2022 – Jul 2024

- Built analytics APIs and optimized Redis caching and MySQL indexing, cutting load times ~25% and query time ~30%.
- Automated operational reporting workflows, reducing manual support overhead ~15%.

SELECTED PROJECTS

Suture — An AI-native operations layer for a cardiology practice: the full referral-to-outreach loop, built to be audited.

Python, FastAPI, PostgreSQL, pgvector, LLM extraction, Hybrid RAG · github.com/Shaurayagulati/suture-ai

rag-verdict — Pytest for RAG agents: an open-source framework that tests what your agent does, not just how its answers score.

Python, pytest, LLM-as-judge, CI, Open source (MIT) · github.com/Shauryagulati/ragverdict

Regulatory RAG at CMU — A hybrid retrieval system over 1,000+ regulatory documents. 94% accuracy with grounded citations.

E5 embeddings, ChromaDB, BM25, BGE reranking, LangChain, RAGAS · github.com/Shauryagulati/REG-RAG

This site — a terminal with a built-in RAG agent over my work, running at \$0/month.

shauryagulati.me — type "shaurya" in its terminal.

PUBLICATIONS

- Enhancing Sentiment Analysis in Short Texts with POS-Embedded LSTM Models
- Paint/Writing Application through Webcam using MediaPipe and OpenCV
- Virtual Assistant with Sentiment Analysis
- Can Artificial Intelligence Work Wonders in Healthcare?

SKILLS

AI / LLM: Agents, RAG, Claude Agent SDK, OpenAI API, LangChain, MCP, LLM-as-judge, Evals, Prompt engineering & caching

RETRIEVAL: pgvector, FAISS, Hybrid search, BM25, RRF fusion, BGE cross-encoder reranking, Query decomposition

BACKEND: Python, FastAPI, PostgreSQL, Redis, SQLAlchemy, Alembic, asyncio, SSE streaming, SQL / MySQL

OBSERVABILITY / INFRA: Langfuse, OpenTelemetry, Docker, GitHub Actions, Git